

# Chem 277B Spring 2024 Tutorial 6

---

## Outline

- Principal Component Analysis (PCA) in sklearn
- Dropout and L2 regularization in PyTorch
- Dataset and DataLoader in PyTorch

## MNIST Dataset

```
In [ ]: import pickle
import torch

def load_dataset(path):
    with open(path, 'rb') as f:
        train_data, test_data = pickle.load(f)


    X_train = torch.tensor(train_data[0], dtype=torch.float)
    y_train = torch.tensor(train_data[1], dtype=torch.long)
    X_test = torch.tensor(test_data[0], dtype=torch.float)
    y_test = torch.tensor(test_data[1], dtype=torch.long)
    return X_train, y_train, X_test, y_test

X_train, y_train, X_test, y_test = load_dataset("mnist.pkl")
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: torch.Size([60000, 32, 32])
X_test shape: torch.Size([10000, 32, 32])
y_train shape: torch.Size([60000])
y_test shape: torch.Size([10000])
```

```
In [ ]: import matplotlib.pyplot as plt

fig, axes = plt.subplots(3, 3, figsize=(5, 5))
for i, ax in enumerate(axes.flatten()):
    ax.imshow(X_train[i], cmap='gray')
```



## Principal Component Analysis (PCA)

```
In [ ]: # Flatten the inputs & normalization
X_train = X_train.reshape(X_train.shape[0], -1) / torch.max(X_train)
X_test = X_test.reshape(X_test.shape[0], -1) / torch.max(X_test)
print(X_train.shape)
```

```
torch.Size([60000, 1024])
```

```
In [ ]: from sklearn.decomposition import PCA

# keeping specific number of features
pca = PCA(n_components=256)
# fit
pca.fit(X_train)
# transform
X_train_pca = pca.transform(X_train)
X_test_pca = pca.transform(X_test)
print(X_train_pca.shape, X_test_pca.shape)
```

```
(60000, 256) (10000, 256)
```

```
In [ ]: # keeping amount of variance
pca = PCA(n_components=0.99)
# fit
pca.fit(X_train)
# transform
X_train_pca = pca.transform(X_train)
```

```
X_test_pca = pca.transform(X_test)
print(X_train_pca.shape, X_test_pca.shape)
```

```
(60000, 331) (10000, 331)
```

```
In [ ]: X_train_pca = torch.tensor(X_train_pca, dtype=torch.float)
X_test_pca = torch.tensor(X_test_pca, dtype=torch.float)
```

```
In [ ]: print(X_train_pca.shape, X_test_pca.shape)

torch.Size([60000, 331]) torch.Size([10000, 331])
```

```
In [ ]: sum(pca.explained_variance_ratio_)
```

```
Out[ ]: 0.990012942492929
```

## Dropout & L2

Use `nn.Dropout` layer:

During training, randomly zeroes some of the elements of the input tensor with probability  $p$ .

```
In [ ]: import torch.nn as nn

class NetDropout(nn.Module):
    def __init__(self):
        super().__init__()
        self.layers = nn.Sequential(
            nn.Linear(331, 100),
            nn.Dropout(p=0.1),
            nn.Sigmoid(),
            nn.Linear(100, 10),
            nn.Dropout(p=0.1),
            nn.Sigmoid()
        )

    def forward(self, x):
        return self.layers(x)
```

Set `weight_decay` to use L2 regularization.

$$\text{Loss}_{L2} = \text{Loss} + \lambda \sum \theta_i^2$$

```
In [ ]: model = NetDropout()
optimizer = torch.optim.Adam(model.parameters(), 1e-3, weight_decay=1e-5)
```

## Dataset & DataLoader

```
In [ ]: from torch.utils.data import Dataset, DataLoader
```

```
class MnistDataset(Dataset):
    def __init__(self, X, y):
        self.X = X
        self.y = y

    def __len__(self):
        return len(self.y)

    def __getitem__(self, idx):
        return self.X[idx], self.y[idx]
```

```
In [ ]: train_data = MnistDataset(X_train_pca, y_train)
test_data = MnistDataset(X_test_pca, y_test)
train_loader = DataLoader(train_data, batch_size=128, shuffle=True)
test_loader = DataLoader(test_data, batch_size=128, shuffle=True)
```

```
In [ ]: for X_batch, y_batch in train_loader:
    print(X_batch.shape, y_batch.shape)
    break
```

```
torch.Size([128, 331]) torch.Size([128])
```

## Trainer

```
In [ ]: import numpy as np
from tqdm import tqdm

class Trainer:

    def __init__(self, model, opt_method, learning_rate, batch_size, epoch,
                 self.model = model

        if opt_method == "adam":
            self.optimizer = torch.optim.Adam(model.parameters(), learning_r
        else:
            raise NotImplementedError("This optimization is not supported")

        self.epoch = epoch
        self.batch_size = batch_size

    def train(self, train_data, val_data, early_stop=True, verbose=True, dra
        train_loader = DataLoader(train_data, batch_size=self.batch_size, sh

        train_loss_list, train_acc_list = [], []
        val_loss_list, val_acc_list = [], []
        weights = self.model.state_dict()
        lowest_val_loss = np.inf
        loss_func = nn.CrossEntropyLoss()
        for n in tqdm(range(self.epoch), leave=False):
            # enable train mode
            self.model.train()
```

```

epoch_loss, epoch_acc = 0.0, 0.0
for X_batch, y_batch in train_loader:
    # batch_importance is the ratio of batch_size
    batch_importance = y_batch.shape[0]/len(train_data)
    y_pred = self.model(X_batch)
    batch_loss = loss_func(y_pred, y_batch)

    self.optimizer.zero_grad()
    batch_loss.backward()
    self.optimizer.step()

    epoch_loss += batch_loss.detach().cpu().item() * batch_importance
    batch_acc = torch.sum(torch.argmax(y_pred, axis=1) == y_batch, axis=0).item()
    epoch_acc += batch_acc.detach().cpu().item() * batch_importance
train_loss_list.append(epoch_loss)
train_acc_list.append(epoch_acc)
val_loss, val_acc = self.evaluate(val_data)
val_loss_list.append(val_loss)
val_acc_list.append(val_acc)

if early_stop:
    if val_loss < lowest_val_loss:
        lowest_val_loss = val_loss
        weights = self.model.state_dict()

if draw_curve:
    x_axis = np.arange(self.epoch)
    fig, axes = plt.subplots(1, 2, figsize=(10, 4))
    axes[0].plot(x_axis, train_loss_list, label="Train")
    axes[0].plot(x_axis, val_loss_list, label="Validation")
    axes[0].set_title("Loss")
    axes[0].legend()
    axes[1].plot(x_axis, train_acc_list, label='Train')
    axes[1].plot(x_axis, val_acc_list, label='Validation')
    axes[1].set_title("Accuracy")
    axes[1].legend()

if early_stop:
    self.model.load_state_dict(weights)

return {
    "train_loss_list": train_loss_list,
    "train_acc_list": train_acc_list,
    "val_loss_list": val_loss_list,
    "val_acc_list": val_acc_list,
}

def evaluate(self, data, print_acc=False):
    # enable evaluation mode
    self.model.eval()
    loader = DataLoader(data, batch_size=self.batch_size, shuffle=True)
    loss_func = nn.CrossEntropyLoss()
    acc, loss = 0.0, 0.0
    for X_batch, y_batch in loader:
        with torch.no_grad():
            batch_importance = y_batch.shape[0]/len(train_data)

```

```
        y_pred = self.model(X_batch)
        batch_loss = loss_func(y_pred, y_batch)
        batch_acc = torch.sum(torch.argmax(y_pred, axis=1) == y_batch, axis=0)
        acc += batch_acc.detach().cpu().item() * batch_importance
        loss += batch_loss.detach().cpu().item() * batch_importance
    if print_acc:
        print(f"Accuracy: {acc:.3f}")
    return loss, acc
```

```
In [ ]: trainer = Trainer(model, "adam", 1e-3, 128, 50, 1e-5)
        trainer.train(train_data, test_data)
```



```
Out[ ]: {'train_loss_list': [1.6345654050827023,  
1.6256413174947117,  
1.619539224243165,  
1.6178119778315228,  
1.6122155399958304,  
1.6093794220606485,  
1.608482451057434,  
1.6063486288070667,  
1.6038572382609038,  
1.6014241783142082,  
1.6007745227177927,  
1.599812180264789,  
1.597846146583558,  
1.5972427301406853,  
1.5956069282531724,  
1.5947243247985832,  
1.5936766189575198,  
1.593202598508199,  
1.5918138434728,  
1.5926415494918846,  
1.5917330230712887,  
1.5907306164423625,  
1.5894648267745985,  
1.5896575952529912,  
1.5875960771560667,  
1.58758466536204,  
1.5862048535664877,  
1.586286747423808,  
1.5872448308308928,  
1.585802490043641,  
1.58577879892985,  
1.5855446199417098,  
1.5844957898457848,  
1.5848428560892744,  
1.5839182889938355,  
1.5830990733464563,  
1.5834927261988339,  
1.5828149834314986,  
1.5822036504745487,  
1.5814984322230017,  
1.58175454451243,  
1.5815143378575645,  
1.5822904492060332,  
1.5805196723302206,  
1.5810784157435107,  
1.5812023804346726,  
1.5794532040913887,  
1.5800079744338982,  
1.5792736788431794,  
1.5789225581487012],  
'train_acc_list': [0.8782166666984605,  
0.8854500000000044,  
0.8912833333333375,  
0.89041666669846,  
0.8952833333651263,  
0.8975166666666704,
```

```
0.90000000000000036,  
0.9019833333015477,  
0.9049166666348809,  
0.9064166666984592,  
0.9075166666666697,  
0.9104333333651254,  
0.912983333301547,  
0.9121833333015469,  
0.9136833333015468,  
0.9153333333651251,  
0.9166333333333356,  
0.91691666663488,  
0.9181666666348799,  
0.9188333333333355,  
0.9206000000000022,  
0.9186333333333355,  
0.920966666698458,  
0.9221666666666686,  
0.922249999968213,  
0.924400000031791,  
0.9245166666984577,  
0.9244666666666684,  
0.9242666666666685,  
0.9261833333651242,  
0.9251666666984576,  
0.9269000000000017,  
0.9269833333651242,  
0.9297333333333346,  
0.9265166666666683,  
0.9292500000000012,  
0.9295500000000015,  
0.9300333333333348,  
0.9310166666666677,  
0.9311499999682122,  
0.9327333333333343,  
0.9309333333333345,  
0.9315333333333344,  
0.9333333333651236,  
0.9324833333333344,  
0.9327666666666676,  
0.9350000000000008,  
0.9331666666348787,  
0.9347000000000001,  
0.9350666666666676],  
'val_loss_list': [0.2588511471748352,  
0.25773222007751473,  
0.25691719643274935,  
0.2563522413253785,  
0.2558027157147726,  
0.25547989222208656,  
0.2551308627128602,  
0.2547252587318421,  
0.25442750892639165,  
0.25422911443710333,  
0.2540464429219565,  
0.2537937870661418,
```



0.2536431848526002,  
0.2534067885080974,  
0.25320728416442884,  
0.25312399593989054,  
0.25289309854507447,  
0.2528072364171346,  
0.25273120349248246,  
0.25260152072906494,  
0.25241151371002196,  
0.252415691057841,  
0.2522326436996459,  
0.25216817359924315,  
0.2520019748369852,  
0.25197507321039825,  
0.2518450228055318,  
0.25169793780644745,  
0.2516503425916036,  
0.25158187856674197,  
0.251466447766622,  
0.2513963534673056,  
0.25126730022430416,  
0.2512035391171774,  
0.2511685480117798,  
0.2510992826461793,  
0.251028265062968,  
0.25103135042190555,  
0.250923332309723,  
0.2508750089009603,  
0.25078297662734994,  
0.25076765464146933,  
0.2507223012606303,  
0.2506392167091371,  
0.2505964373906454,  
0.25054782549540205,  
0.25052587922414143,  
0.2505193968772888,  
0.25040930995941174,  
0.25038921445210777],  
'val\_acc\_list': [0.15325,  
0.15410000000000004,  
0.15448333333333334,  
0.15483333333333338,  
0.15510000000000002,  
0.15546666666666667,  
0.15543333333333334,  
0.15590000000000004,  
0.15626666666666667,  
0.156316666666666674,  
0.15641666666666665,  
0.15664999999999996,  
0.15681666666666666,  
0.157016666666666658,  
0.15715000000000007,  
0.15733333333333333,  
0.15738333333333332,  
0.15756666666666666,

```
0.15773333333333336,  
0.15776666666666667,  
0.15793333333333331,  
0.15803333333333333,  
0.15808333333333333,  
0.15818333333333333,  
0.15848333333333333,  
0.15844999999999995,  
0.15871666666666666,  
0.15891666666666665,  
0.15891666666666666,  
0.15886666666666663,  
0.15923333333333333,  
0.15909999999999996,  
0.15939999999999999,  
0.15936666666666663,  
0.15944999999999995,  
0.15958333333333322,  
0.15973333333333328,  
0.15969999999999995,  
0.15981666666666666,  
0.15978333333333328,  
0.15981666666666663,  
0.16001666666666658,  
0.15984999999999996,  
0.15981666666666666,  
0.15991666666666666,  
0.15998333333333333,  
0.16016666666666666,  
0.16008333333333322,  
0.16013333333333333,  
0.16018333333333332] }
```



In [ ]: