

# Chem 277B Spring 2024 Tutorial 7

---

## Outline

- Convolutional Neural Network (CNN):
  - Hyperparameters in CNN: channels, padding, stride, dilation
  - Pooling
  - CNN in PyTorch
- Residual Network
- Batch Normalization

## HW6 - Helper function

You can use the following decorator to report time:

```
In [ ]: import time

def timeit(f):

    def timed(*args, **kw):

        ts = time.time()
        result = f(*args, **kw)
        te = time.time()

        print(f'func:{f.__name__} took: {te-ts:.4f} sec')
        return result

    return timed

@timeit
def sleep(sec):
    return time.sleep(sec)

sleep(0.1)
```

func:sleep took: 0.1050 sec

```
In [ ]: class Trainer:

    def __init__(self, model, opt_method, learning_rate, batch_size, epoch,
                self.model = model

    if opt_method == "adam":
        self.optimizer = torch.optim.Adam(model.parameters(), learning_r
```

```

else:
    raise NotImplementedError("This optimization is not supported")

self.epoch = epoch
self.batch_size = batch_size

@timeit
def train(self, train_data, val_data, early_stop=True, verbose=True, draw_curve=True,
          train_loader = DataLoader(train_data, batch_size=self.batch_size, shuffle=True)):

    train_loss_list, train_acc_list = [], []
    val_loss_list, val_acc_list = [], []
    weights = self.model.state_dict()
    lowest_val_loss = np.inf
    loss_func = nn.CrossEntropyLoss()
    for n in tqdm(range(self.epoch), leave=False):
        # enable train mode
        self.model.train()
        epoch_loss, epoch_acc = 0.0, 0.0
        for X_batch, y_batch in train_loader:
            # batch_importance is the ratio of batch_size
            batch_importance = y_batch.shape[0]/len(train_data)
            y_pred = self.model(X_batch)
            batch_loss = loss_func(y_pred, y_batch)

            self.optimizer.zero_grad()
            batch_loss.backward()
            self.optimizer.step()

            epoch_loss += batch_loss.detach().cpu().item() * batch_importance
            batch_acc = torch.sum(torch.argmax(y_pred, axis=1) == y_batch, dim=0).item()
            epoch_acc += batch_acc.detach().cpu().item() * batch_importance

        # train_loss_list.append(epoch_loss)
        # train_acc_list.append(epoch_acc)
        # previous way to report might get low acc due to dropout
        train_loss, train_acc = self.evaluate(train_data)

        val_loss, val_acc = self.evaluate(val_data)
        val_loss_list.append(val_loss)
        val_acc_list.append(val_acc)

        if early_stop:
            if val_loss < lowest_val_loss:
                lowest_val_loss = val_loss
                weights = self.model.state_dict()

    if draw_curve:
        x_axis = np.arange(self.epoch)
        fig, axes = plt.subplots(1, 2, figsize=(10, 4))
        axes[0].plot(x_axis, train_loss_list, label="Train")
        axes[0].plot(x_axis, val_loss_list, label="Validation")
        axes[0].set_title("Loss")
        axes[0].legend()
        axes[1].plot(x_axis, train_acc_list, label='Train')

```

```

axes[1].plot(x_axis, val_acc_list, label='Validation')
axes[1].set_title("Accuracy")
axes[1].legend()

if early_stop:
    self.model.load_state_dict(weights)


return {
    "train_loss_list": train_loss_list,
    "train_acc_list": train_acc_list,
    "val_loss_list": val_loss_list,
    "val_acc_list": val_acc_list,
}

def evaluate(self, data, print_acc=False):
    # enable evaluation mode
    self.model.eval()
    loader = DataLoader(data, batch_size=self.batch_size, shuffle=True)
    loss_func = nn.CrossEntropyLoss()
    acc, loss = 0.0, 0.0
    for X_batch, y_batch in loader:
        with torch.no_grad():
            batch_importance = y_batch.shape[0]/len(data)
            y_pred = self.model(X_batch)
            batch_loss = loss_func(y_pred, y_batch)
            batch_acc = torch.sum(torch.argmax(y_pred, axis=1) == y_batch, axis=0).item()
            acc += batch_acc.detach().cpu().item() * batch_importance
            loss += batch_loss.detach().cpu().item() * batch_importance
    if print_acc:
        print(f"Accuracy: {acc:.3f}")
    return loss, acc


```

## Convolutional Neural Network (CNN)

### CNN general architecture



### Convolution Filters help extract features



## Calculating convolution output shape


Here is a [visualization](#) for padding, stride and dilation

$$H_{\text{out}} = \left\lfloor \frac{H_{\text{in}} + 2 \times \text{padding} - \text{dilation} \times (\text{kernel size} - 1) - 1}{\text{stride}} + 1 \right\rfloor$$

```
In [ ]: import pickle
import torch
import torch.nn as nn
```

```
In [ ]: # init a Conv2d layer
conv = nn.Conv2d(1, 1, kernel_size=(2,2))
conv
```

```
Out[ ]: Conv2d(1, 1, kernel_size=(2, 2), stride=(1, 1))
```



```
In [ ]: # init a MaxPool layer
max_pool = nn.MaxPool2d(2)
max_pool
```

```
Out[ ]: MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
```

```
In [ ]: # init a Average Pool layer
avg_pool = nn.AvgPool2d(2)
```

```
avg_pool
```

```
Out[ ]: AvgPool2d(kernel_size=2, stride=2, padding=0)
```

```
In [ ]: def out_dim(in_dim, kernel_size, padding, stride, dilation):
        return (in_dim + 2 * padding - dilation * (kernel_size - 1) - 1) // stride

# data shape: (N, C, W, H)
data = torch.rand(1, 1, 28, 28)
conv(data)
```

```
Out[ ]: tensor([[[[-0.4292]]]], grad_fn=<ConvolutionBackward0>)
```

## LeNet architecture

LeCun, Y.; Bottou, L.; Bengio, Y. & Haffner, P. (1998). Gradient-based learning applied to document recognition. Proceedings of the IEEE. 86(11): 2278 - 2324.

Layer No.	Layer type	#channels/#features	Kernel size	Stride	Activation
1	2D Convolution	6	5	1	tanh
2	Average pooling	6	2	2	\
3	2D Convolution	16	5	1	tanh
4	Average pooling	16	2	2	\
5	2D Convolution	120	5	1	tanh
6	Flatten				\
7	Fully connected	84			tanh
8	Fully connected	10			softmax

```
In [ ]: def load_dataset(path):
        with open(path, 'rb') as f:
            train_data, test_data = pickle.load(f)

        X_train = torch.tensor(train_data[0], dtype=torch.float).unsqueeze(1)
        y_train = torch.tensor(train_data[1], dtype=torch.long)
        X_test = torch.tensor(test_data[0], dtype=torch.float).unsqueeze(1)
        y_test = torch.tensor(test_data[1], dtype=torch.long)
        return X_train, y_train, X_test, y_test

X_train, y_train, X_test, y_test = load_dataset("mnist.pkl")
```

```
In [ ]: class LeNet(nn.Module):
        def __init__(self, in_channels=1):
            super().__init__()
            self.conv = nn.ModuleList([
                nn.Conv2d(in_channels, 6, kernel_size=5, stride=1),
                nn.Conv2d(6, 16, kernel_size=5, stride=1),
```

```

        nn.Conv2d(16, 120, kernel_size=5, stride=1)
    ])
    self.pool = nn.AvgPool2d(2)
    self.activation = nn.Tanh()
    self.fc = nn.ModuleList([
        nn.Linear(120, 84),
        nn.Linear(84, 10)
    ])

    def forward(self, x):
        for i in range(2):
            x = self.pool(self.activation(self.conv[i](x)))
        x = nn.Flatten()(self.activation(self.conv[2](x)))
        x = self.activation(self.fc[0](x))
        x = nn.Softmax(dim=-1)(self.fc[1](x))
        return x

net = LeNet()
net

```

```

Out [ ]: LeNet(
  (conv): ModuleList(
    (0): Conv2d(1, 6, kernel_size=(5, 5), stride=(1, 1))
    (1): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
    (2): Conv2d(16, 120, kernel_size=(5, 5), stride=(1, 1))
  )
  (pool): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (activation): Tanh()
  (fc): ModuleList(
    (0): Linear(in_features=120, out_features=84, bias=True)
    (1): Linear(in_features=84, out_features=10, bias=True)
  )
)

```

```

In [ ]: # Use torchsummary to print the architecture
# ! pip install torchsummary
from torchsummary import summary

s = summary(net, (1, 32, 32))

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 6, 28, 28]	156
Tanh-2	[-1, 6, 28, 28]	0
AvgPool2d-3	[-1, 6, 14, 14]	0
Conv2d-4	[-1, 16, 10, 10]	2,416
Tanh-5	[-1, 16, 10, 10]	0
AvgPool2d-6	[-1, 16, 5, 5]	0
Conv2d-7	[-1, 120, 1, 1]	48,120
Tanh-8	[-1, 120, 1, 1]	0
Linear-9	[-1, 84]	10,164
Tanh-10	[-1, 84]	0
Linear-11	[-1, 10]	850

Total params: 61,706

Trainable params: 61,706

Non-trainable params: 0

Input size (MB): 0.00

Forward/backward pass size (MB): 0.11

Params size (MB): 0.24

Estimated Total Size (MB): 0.35

In [ ]: `net(X_train[:10])`

```

Out[ ]: tensor([[0.0852, 0.0939, 0.0924, 0.0957, 0.0943, 0.1224, 0.0921, 0.0951, 0.
1162,
          0.1127],
          [0.0886, 0.0982, 0.0959, 0.0981, 0.0989, 0.1172, 0.0908, 0.0948, 0.
1115,
          0.1060],
          [0.0857, 0.0923, 0.0987, 0.0914, 0.0947, 0.1241, 0.1011, 0.0948, 0.
1085,
          0.1088],
          [0.0866, 0.0932, 0.0929, 0.1049, 0.0901, 0.1173, 0.1016, 0.0951, 0.
1060,
          0.1124],
          [0.0814, 0.0931, 0.0989, 0.1023, 0.0948, 0.1180, 0.0936, 0.0964, 0.
1132,
          0.1083],
          [0.0857, 0.0949, 0.0938, 0.1034, 0.0925, 0.1185, 0.0992, 0.0949, 0.
1068,
          0.1102],
          [0.0846, 0.0943, 0.0928, 0.0976, 0.0920, 0.1193, 0.0957, 0.1011, 0.
1092,
          0.1133],
          [0.0861, 0.0948, 0.0940, 0.0980, 0.0971, 0.1200, 0.0924, 0.0957, 0.
1109,
          0.1110],
          [0.0840, 0.0929, 0.0944, 0.0982, 0.0943, 0.1147, 0.1005, 0.1016, 0.
1086,
          0.1107],
          [0.0867, 0.0931, 0.0925, 0.1038, 0.0950, 0.1199, 0.0986, 0.0930, 0.
1104,
          0.1069]], grad_fn=<SoftmaxBackward0>)

```

## Residual Network (ResNet)

An example of residual block:




Figure 2. Residual learning: a building block.

```

In [ ]: class ResBlock(nn.Module):
        def __init__(self, dim):
            super().__init__()

```



```

        self.fc = nn.ModuleList([nn.Linear(dim, dim), nn.Linear(dim, dim)])
        self.activation = nn.ReLU()

    def forward(self, x):
        out = self.activation(self.fc[0](x))
        out = self.fc[1](out)
        out += x
        out = self.activation(out)
        return out

```

```

In [ ]: class LeNetRes(nn.Module):
    def __init__(self, in_channels=1):
        super().__init__()
        self.conv = nn.ModuleList([
            nn.Conv2d(in_channels, 6, kernel_size=5, stride=1),
            nn.Conv2d(6, 16, kernel_size=5, stride=1),
            nn.Conv2d(16, 120, kernel_size=5, stride=1)
        ])
        self.pool = nn.AvgPool2d(2)
        self.activation = nn.Tanh()
        self.fc = nn.ModuleList([
            nn.Linear(120, 120),
            nn.Linear(120, 84),
            nn.Linear(84, 10)
        ])

    def forward(self, x):
        for i in range(2):
            x = self.pool(self.activation(self.conv[i](x)))
        x = nn.Flatten()(self.activation(self.conv[2](x)))
        x = self.activation(x + self.fc[0](x))
        x = self.activation(self.fc[0](x))
        x = nn.Softmax(dim=-1)(self.fc[1](x))
        return x

net = LeNetRes()
net

```

```

Out [ ]: LeNetRes(
  (conv): ModuleList(
    (0): Conv2d(1, 6, kernel_size=(5, 5), stride=(1, 1))
    (1): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
    (2): Conv2d(16, 120, kernel_size=(5, 5), stride=(1, 1))
  )
  (pool): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (activation): Tanh()
  (fc): ModuleList(
    (0): Linear(in_features=120, out_features=120, bias=True)
    (1): Linear(in_features=120, out_features=84, bias=True)
    (2): Linear(in_features=84, out_features=10, bias=True)
  )
)

```

```

In [ ]: s = summary(net, (1, 32, 32))

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 6, 28, 28]	156
Tanh-2	[-1, 6, 28, 28]	0
AvgPool2d-3	[-1, 6, 14, 14]	0
Conv2d-4	[-1, 16, 10, 10]	2,416
Tanh-5	[-1, 16, 10, 10]	0
AvgPool2d-6	[-1, 16, 5, 5]	0
Conv2d-7	[-1, 120, 1, 1]	48,120
Tanh-8	[-1, 120, 1, 1]	0
Linear-9	[-1, 120]	14,520
Tanh-10	[-1, 120]	0
Linear-11	[-1, 120]	14,520
Tanh-12	[-1, 120]	0
Linear-13	[-1, 84]	10,164
Total params: 89,896		
Trainable params: 89,896		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.11		
Params size (MB): 0.34		
Estimated Total Size (MB): 0.46		

## Batch Normalization (BN)

For a 4-D input data  $X$  with shape  $(N, C, W, H)$ . For each channel, the data is normalized by:

$$\hat{X}_{ijkl} = \frac{X_{ijkl} - \text{mean}(X_j)}{\sqrt{\text{var}(X_j) + \epsilon}} * \gamma_j + \beta_j$$

where

$$\text{mean}(X_j) = \frac{1}{NWH} \sum_i^N \sum_k^W \sum_l^H X_{ikl}$$

$$\text{var}(X_j) = \frac{1}{NWH} \sum_i^N \sum_k^W \sum_l^H (X_{ikl} - \text{mean}(X_j))^2$$

$\epsilon$  is a small number (say,  $10^{-5}$ ) to avoid numerical instability.  $\gamma, \beta$  are learnable parameters

```
In [ ]: batch_norm = nn.BatchNorm2d(120)
batch_norm
```

```
Out[ ]: BatchNorm2d(120, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
```

```
In [ ]: class LeNetResNorm(nn.Module):
    def __init__(self, in_channels=1):
        super().__init__()
        self.conv = nn.ModuleList([
            nn.Conv2d(in_channels, 6, kernel_size=5, stride=1),
            nn.Conv2d(6, 16, kernel_size=5, stride=1),
            nn.Conv2d(16, 120, kernel_size=5, stride=1)
        ])
        self.bn = nn.ModuleList([
            nn.BatchNorm2d(6),
            nn.BatchNorm2d(16),
        ])
        self.pool = nn.AvgPool2d(2)
        self.activation = nn.Tanh()
        self.fc = nn.ModuleList([
            nn.Linear(120, 120),
            nn.Linear(120, 84),
            nn.Linear(84, 10)
        ])

    def forward(self, x):
        for i in range(2):
            x = self.bn[i](self.pool(self.activation(self.conv[i](x))))
        x = nn.Flatten()(self.activation(self.conv[2](x)))
        x = self.activation(x + self.fc[0](x))
        x = self.activation(self.fc[0](x))
        x = nn.Softmax(dim=-1)(self.fc[1](x))
        return x

net = LeNetResNorm()
net
```

```
Out[ ]: LeNetResNorm(
  (conv): ModuleList(
    (0): Conv2d(1, 6, kernel_size=(5, 5), stride=(1, 1))
    (1): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
    (2): Conv2d(16, 120, kernel_size=(5, 5), stride=(1, 1))
  )
  (bn): ModuleList(
    (0): BatchNorm2d(6, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (pool): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (activation): Tanh()
  (fc): ModuleList(
    (0): Linear(in_features=120, out_features=120, bias=True)
    (1): Linear(in_features=120, out_features=84, bias=True)
    (2): Linear(in_features=84, out_features=10, bias=True)
  )
)
```

In [ ]: